Adaptive Machine Learning Models for Streaming Financial Data Using Meta-Heuristic Optimization

Siva Sabarivel KM, VinothKumar A

Abstract

Machine learning models deployed in financial environments such as stock trading systems, algorithmic trading engines, fraud detection pipelines, and market risk analysis tools are traditionally trained on historical data under the assumption that future statistical properties remain stable. However, financial data streams exhibit high non-stationarity, with evolving patterns driven by economic cycles, regulatory changes, geopolitical events, and behavioral shifts among market participants. These dynamic patterns introduce concept drift, which progressively degrades the performance of static machine learning models. Recent research has explored meta-heuristic optimization techniques as adaptive mechanisms enabling autonomous model responses to distributional changes. Meta-heuristics provide flexible, population-based search strategies that dynamically adjust model configurations without gradient information. This literature review synthesizes ten research studies published between 2021 and 2025 investigating adaptive machine learning strategies for streaming data with emphasis on financial applications. The review analyzes how different meta-heuristic algorithms integrate with drift detection mechanisms, feature adaptation processes, and model optimization workflows, evaluates their reported effectiveness, and identifies persistent challenges and future research directions.

Introduction

Machine learning adoption within the financial sector has accelerated over the past decade, driven by large-scale data stream availability and computational advances. Applications including stock price prediction, credit scoring, fraud detection, and automated trading rely heavily on predictive models processing continuous financial data streams. These data streams are inherently dynamic, characterized by frequent distributional changes caused by market volatility, policy reforms, technological innovation, and unforeseen global events.

Conventional machine learning models are typically trained offline using historical datasets and deployed under the assumption that future data follows similar statistical patterns. When this assumption fails, predictive performance deteriorates—a degradation commonly attributed to concept drift, a well-documented challenge in data stream mining. Concept drift represents a fundamental obstacle to long-term deployment of reliable machine learning systems in real-time financial environments. The financial sector presents unique challenges due to the rapid pace of market changes, where concept drift can manifest within hours during crisis periods such as market crashes or sudden regulatory announcements.

Adaptive learning strategies have been proposed enabling models to update continuously as new data arrives. Among these strategies, meta-heuristic optimization has emerged as a promising approach due to its ability to efficiently search complex and evolving solution spaces without requiring gradient information or complete model retraining. Meta-heuristic algorithms such as Genetic Algorithms, Particle Swarm Optimization, and evolutionary computation techniques offer population-based search strategies that can explore multiple solution candidates simultaneously,

making them particularly suited for dynamic optimization in non-stationary environments. This review examines how meta-heuristic optimization supports adaptive machine learning models under concept drift, focusing on financial data streams. The review addresses three research questions: (1) how meta-heuristic algorithms integrate into drift-handling frameworks, (2) which adaptation strategies prove most effective in financial streaming contexts, and (3) what limitations and research gaps remain unresolved in current research.

Literature Review

Paper 1: Riess (2022)

Riess presented a comprehensive survey examining meta-heuristic optimization roles in concept drift adaptation, specifically analyzing Full Model Selection (FMS) frameworks that integrate Feature Selection, Model Selection, and Hyperparameter Optimization into unified adaptive systems. The study demonstrated that static machine learning models prove unsuitable for non-stationary environments due to their assumption of fixed data distributions. Population-based meta-heuristic algorithms such as Genetic Algorithms and Particle Swarm Optimization emerged as particularly effective for dynamic optimization tasks, enabling continuous reconfiguration of model parameters and structures as data evolves while reducing dependence on repeated full retraining cycles. The research documented a ten-year transition in the field from isolated AutoML tasks toward integrated FMS frameworks, establishing meta-heuristic optimization as a practical mechanism for automating model adaptation in streaming environments. The study analyzed over 100 publications spanning 2010-2022, identifying key trends including increased adoption of ensemble-based approaches combining multiple meta-heuristic algorithms for improved robustness. However, the survey also identified critical gaps in evaluation transparency, noting insufficient documentation of drift type testing and inconsistent benchmarking methodologies across studies. Riess emphasized that future research must develop standardized evaluation protocols that systematically test adaptive systems across all drift types—sudden, gradual, incremental, and recurring—rather than focusing on single drift scenarios.

Paper 2: Agrahari (2022)

Agrahari and Singh conducted a systematic investigation of statistical methods for detecting concept drift in data stream mining environments. The study examined hypothesis testing approaches including Kolmogorov-Smirnov tests, Chi-squared tests, and t-tests, alongside sliding window techniques such as ADWIN (Adaptive Windowing) and fixed-size windows, and prediction error monitoring mechanisms including Page-Hinkley tests and CUSUM (Cumulative Sum) control charts. Empirical results demonstrated that sustained increases in error rates reliably signal significant distributional changes in streaming data. The research compared detection delay, false positive rates, and computational overhead across fifteen different drift detection methods using synthetic datasets with controlled drift characteristics. Results showed that ADWIN achieved the best balance between detection speed (average delay of 47 samples) and false positive rate (3.2%), while simpler methods like fixed sliding windows exhibited higher false positive rates (11-15%) but lower computational costs. While the research focused primarily on drift detection rather than adaptation mechanisms, it established the fundamental importance of timely detection as a prerequisite for adaptive learning systems. The findings provide essential groundwork for optimization-based adaptation strategies that respond once drift is detected, emphasizing that

detection delays directly impact the effectiveness of subsequent adaptation procedures regardless of optimization sophistication.

Paper 3: Yelleti et al. (2024)

Yelleti, Ravi, and Radha Krishna proposed a groundbreaking online feature subset selection framework specifically designed for streaming data under big data paradigms, addressing the critical challenge of feature selection when both data samples and features arrive at high velocity in time-dependent streams. The research introduced three key innovations: (1) a parallel hybrid Particle Swarm Optimization wrapper combining Bare-Bones PSO guided by logistic distribution (BBPSO-L) with Threshold-based Neighborhood Search heuristic (TNS) to achieve superior exploitation capability and avoid local optima entrapment, (2) a scalable implementation leveraging Apache Spark distributed computing framework enabling processing of high-velocity streams in real-time, and (3) two vigilance tests for novelty detection that determine when wrapper retraining is necessary, significantly reducing computational overhead. The proposed BBPSO-L+TNS hybrid demonstrated substantial improvements over baseline evolutionary algorithms, achieving 94.7% classification accuracy on cybersecurity datasets while reducing feature dimensionality by 67% and maintaining sub-second response times for feature selection decisions. Traditional PSO-based methods achieved only 87.3% accuracy with 42% dimensionality reduction, highlighting the effectiveness of the hybrid approach. The framework addresses a critical gap in streaming machine learning where features themselves evolve over time—a scenario common in financial applications where new technical indicators, sentiment signals, or market microstructure features continuously emerge. Experimental validation on NSL-KDD intrusion detection and bioinformatics datasets confirmed scalability up to 10 million samples with 5,000 features, demonstrating practical applicability for high-frequency financial trading systems. The vigilance mechanism reduced unnecessary retraining by 73%, achieving a balance between adaptation responsiveness and computational efficiency essential for real-world deployment.

Paper 4: Mahmood and Fatima (2021)

Mahmood and Fatima conducted a systematic literature review analyzing concept drift in streaming data environments, examining a wide range of drift detection and adaptation techniques across 127 peer-reviewed publications from 2010-2021. The review systematically categorized approaches into three primary strategies: blind adaptation methods that continuously update models without explicit drift detection, informed adaptation methods triggered by detected changes, and hybrid approaches combining both strategies. Key limitations identified related to scalability challenges when processing millions of samples per second, computational efficiency bottlenecks particularly in resource-constrained edge computing scenarios, and the trade-off between adaptation speed and stability. The review emphasized that numerous existing approaches rely on predefined rules or fixed thresholds—such as retraining when error exceeds 15% or updating every 1000 samples—which severely limits their flexibility in dynamic settings where drift characteristics vary unpredictably across different market regimes. Mahmood and Fatima argued that intelligent optimization techniques, particularly meta-heuristics, could substantially improve adaptability by enabling data-driven decision-making during model updates rather than depending on manually tuned parameters that may become obsolete as data patterns evolve. The systematic analysis revealed that only 23% of reviewed studies explicitly addressed computational efficiency alongside accuracy, and fewer than 10% evaluated performance across multiple drift types simultaneously. This comprehensive review reinforced the potential value of meta-heuristic

approaches that adaptively determine optimal configuration parameters based on observed performance rather than static rules, while also highlighting the critical need for multi-objective optimization frameworks that balance accuracy, latency, and computational resource consumption.

Paper 5: Rabash et al. (2023)

Rabash, Nazri, and Shapii expanded the concept drift discussion by introducing feature drift as a distinct phenomenon where the relevance of input features changes over time, independent of target concept shifts. Their comprehensive survey analyzed 89 publications focusing specifically on feature-level changes in streaming environments, demonstrating that feature drift can significantly degrade predictive performance, particularly in high-dimensional financial datasets where feature importance varies across market regimes. The distinction between concept drift (changes in the relationship between features and target) and feature drift (changes in feature relevance or availability) proved critical for developing effective adaptation strategies. For instance, during market transitions from bull to bear conditions, momentum-based technical indicators may lose predictive power while volatility-based indicators gain importance, even though the underlying prediction task remains unchanged. The research showed that evolutionary feature selection methods can dynamically identify relevant feature subsets, thereby improving long-term model stability without requiring complete model reconstruction. Specifically, genetic algorithm-based feature selection reduced dimensionality by 40-60% while maintaining or improving prediction accuracy by 3-8 percentage points compared to using all available features. The survey provided detailed taxonomies of feature drift types including feature disappearance (sensors failing or data sources becoming unavailable), feature emergence (new data sources becoming available), and feature evolution (gradual changes in feature distributions or correlations). This work reinforced the critical importance of optimization-driven feature adaptation in streaming learning systems, demonstrating that meta-heuristic algorithms capable of continuous feature subset evaluation and selection provide substantial advantages over static feature engineering approaches that assume constant feature relevance. The authors recommended hierarchical adaptation strategies where feature selection operates at a faster timescale than model parameter adaptation, enabling more responsive systems that can quickly adjust to changing feature landscapes without triggering computationally expensive model retraining.

Paper 6: Hinder et al. (2024)

Hinder, Vaquet, and Hammer presented a comprehensive survey on monitoring in evolving environments, focusing specifically on unsupervised drift detection methods that operate without labeled data—a critical requirement for many financial applications where obtaining ground truth labels for real-time predictions proves impractical or impossible. The research systematically evaluated fourteen state-of-the-art drift detection methods including MMD (Maximum Mean Discrepancy), DAWIDD (Drift-Aware Windowing with Distribution Distance), KCpD (Kernel Change Point Detection), and ShapeDD (Shape-based Drift Detection) across multiple dimensions: detection capability for varying drift magnitudes, computational efficiency measured in samples processed per second, false positive rates under stable conditions, and detection delay quantified as samples required to identify drift. Experimental results demonstrated that ShapeDD performed particularly well across drift detection scenarios, achieving 97% detection rate for moderate drifts with only 2.1% false positive rate and average detection delay of 31 samples. The study emphasized that meta-heuristic components within ShapeDD—specifically its adaptive threshold adjustment mechanism using MMD kernel optimization—proved crucial for its superior

performance. For simpler datasets, logistic regression-based detectors achieved comparable performance with lower computational cost (processing 50,000 samples/second vs. 8,000 samples/second for kernel methods), while complex high-dimensional datasets required more sophisticated kernel-based approaches for accurate drift detection. The research provided critical insights for financial streaming applications where false positives trigger unnecessary and expensive model updates while delayed detection allows performance degradation during critical market events. The study recommended hybrid detection strategies combining fast lightweight detectors for initial screening with more sophisticated kernel methods for confirmation, reducing false positives by 67% while maintaining detection delays under 50 samples. This work established meta-heuristic optimization as integral not only to model adaptation but also to the drift detection process itself, with adaptive threshold tuning and kernel parameter optimization significantly impacting detector reliability and responsiveness.

Paper 7: Mansour et al. (2021)

Mansour, Al-Otaibi, Al-Rasheed, Aljuaid, Pustokhina, and Pustokhin introduced the Multi-Objective Metaheuristic Optimization-based Big Data Analytics with Concept Drift Detection (MOMBD-CDD) framework specifically designed for high-dimensional streaming data applications with severe class imbalance and concept drift. The framework addresses two critical challenges simultaneously: class imbalance common in financial applications (fraudulent transactions representing 0.1-1% of total volume) and distributional shifts requiring model adaptation. MOMBD-CDD employs Synthetic Minority Over-sampling Technique (SMOTE) to handle class imbalance, with Glowworm Swarm Optimization (GSO) determining optimal oversampling rates and neighboring point selections. GSO, a nature-inspired meta-heuristic algorithm, simulates the behavior of glowworms attracting each other based on luciferin brightness, enabling efficient exploration of the oversampling parameter space. For drift detection, the framework utilizes Statistical Test of Equal Proportions (STEPD), which monitors classification accuracy across sliding windows and triggers adaptation when statistically significant performance degradation is detected. The classification model employs a bidirectional Long Short-Term Memory (Bi-LSTM) network capable of capturing both forward and backward temporal dependencies in sequential financial data. Crucially, GSO-based hyperparameter tuning optimizes the Bi-LSTM configuration including hidden layer dimensions, dropout rates, learning rates, and batch sizes, explicitly balancing classification accuracy against computational efficiency through a multi-objective fitness function. Experimental results on the NSL-KDD intrusion detection dataset demonstrated significant performance improvements: the framework achieved 94.3% accuracy compared to 89.1% for baseline LSTM without optimization, 91.7% for Random Forest, and 87.4% for Support Vector Machines. More importantly for real-time financial applications, drift recovery time decreased dramatically from approximately 500 samples to 87 samples, representing an 82% reduction in adaptation latency. The multi-objective optimization explicitly addresses the accuracy-efficiency trade-off inherent in real-time data stream processing, producing Pareto-optimal solutions that enable system designers to select appropriate operating points based on application-specific constraints. For high-frequency trading requiring sub-millisecond decisions, the framework can prioritize computational efficiency while accepting minor accuracy trade-offs, whereas fraud detection systems can prioritize accuracy given their tolerance for seconds-scale latency.

Paper 8: Xiang et al. (2023)

Xiang, Zi, Cong, and Wang reviewed concept drift adaptation methods within deep learning frameworks, providing a comprehensive taxonomy categorizing approaches based on discriminant learning (classification-focused), generative learning (distribution modeling), and hybrid architectures combining both paradigms. The study systematically analyzed 73 recent publications (2018-2023) focusing on neural network adaptations for non-stationary environments, demonstrating that deep neural networks can integrate with meta-heuristic optimization to dynamically adjust network parameters and architectures in response to detected drift. The research distinguished between active adaptation modes, which employ explicit drift detection algorithms (ADWIN, DDM, EDDM) to trigger updates, and passive modes, which continuously adjust models as data arrives without formal detection mechanisms through techniques like online gradient descent or momentum-based updates. Analysis showed that optimization-based adaptation improves model robustness under complex drift scenarios, particularly in high-dimensional financial datasets where feature interactions change substantially over time. For instance, convolutional neural networks for price pattern recognition achieved 87% accuracy with periodic retraining versus 91% with genetic algorithm-based architecture search that dynamically adjusted filter sizes and layer depths based on current data characteristics. However, the review also highlighted increased computational overhead as a significant limitation: deep learning models combined with meta-heuristic optimization required 15-50x more computational resources than traditional machine learning approaches. Specifically, neural architecture search using evolutionary algorithms consumed 2-8 GPU-hours per adaptation cycle, compared to 1-5 CPU-minutes for traditional model retraining. This computational burden may exceed latency requirements for high-frequency trading applications where decision times must remain below one millisecond, limiting deep learning with meta-heuristic optimization to medium-frequency applications tolerating seconds-scale adaptation. The review identified promising research directions including transfer learning approaches where meta-heuristics optimize domain adaptation parameters rather than complete architectures, potentially reducing computational overhead by 70-80% while maintaining adaptation effectiveness.

Paper 9: Bousbaa et al. (2023)

Bousbaa, Sanchez-Medina, and Bencharef investigated financial time series forecasting using data stream mining techniques, focusing specifically on EUR/USD exchange rate prediction with comprehensive experimental validation spanning 22 years of FOREX data (2000-2022). The research emphasized the non-stationary nature of financial markets and demonstrated that adaptive learning frameworks substantially outperform static models during volatile market periods, with particularly dramatic improvements during crisis events including the 2008 financial crisis, 2011 European debt crisis, and 2020 COVID-19 pandemic. The proposed framework employs Stochastic Gradient Descent (SGD) optimized with Particle Swarm Optimization (PSO), integrated with adaptive sliding windows whose size varies based on Augmented Dickey-Fuller stationarity tests. During stable market conditions characterized by consistent trends and low volatility, the window expands to capture established patterns spanning up to 15 trading days, enabling the model to learn from sufficient historical context. When stationarity tests detect concept drift through significant changes in time series statistical properties (p-values dropping below 0.05), the system triggers window contraction to single-sample updates for rapid adaptation, minimizing reliance on potentially obsolete historical patterns. The PSO component operates every 60 days to optimize SGD hyperparameters including learning rate, momentum coefficient, and regularization strength, preventing local minima entrapment that often occurs with pure gradient-

based optimization in non-stationary environments. Experimental results demonstrated that adaptive window sizing achieved 82% directional prediction accuracy (correctly predicting price movement direction) and reduced mean squared error by 23% compared to fixed-window approaches using constant 15-day windows. During the 2008 financial crisis period (September-December), the adaptive system maintained 76% accuracy while fixed-window methods degraded to 58% accuracy, demonstrating superior resilience under extreme market conditions. Computational profiling revealed that PSO optimization cycles consumed 37 seconds on average, acceptable for daily trading strategies but prohibitive for intraday high-frequency applications. The research provided compelling empirical evidence that meta-heuristic optimization, while not the central algorithmic focus, significantly enhances continuous model adaptation in financial data streams through intelligent hyperparameter tuning that accounts for evolving market dynamics.

Paper 10: Liu et al. (2023)

Liu, Godahewa, Bandara, and Bergmeir addressed concept drift challenges in global time series forecasting, where a single model trained across multiple related time series must maintain performance as distributional characteristics evolve across different series. This scenario directly applies to financial portfolio management where models must simultaneously forecast returns across hundreds of stocks or currencies, each potentially experiencing independent drift patterns. The research proposed a hybrid approach combining ensemble learning, transfer learning, and evolutionary optimization to handle diverse drift characteristics across series families. The framework maintains a population of specialized forecasting models, each optimized for specific drift patterns (trending, seasonal, volatile, mean-reverting), with genetic algorithms determining optimal model combinations for current data characteristics. Experimental evaluation on financial datasets including daily stock returns for S&P 500 constituents and cryptocurrency prices demonstrated that the adaptive ensemble approach outperformed single global models by 12-18% in forecast accuracy measured by symmetric mean absolute percentage error (sMAPE). During market regime transitions such as the 2022 cryptocurrency crash, the evolutionary optimization successfully reweighted ensemble components, increasing allocation to volatility-aware models from 15% to 61% within 3 days of drift detection, enabling faster recovery to acceptable forecast accuracy. The meta-learning component enabled the system to recognize recurring drift patterns across different time periods and series, automatically retrieving and adapting previously successful model configurations when similar patterns reappear—a capability particularly valuable in financial markets exhibiting cyclical behavior across economic cycles. Computational analysis showed that ensemble model selection required 150-300 milliseconds per forecast horizon, acceptable for daily or hourly forecasting but potentially limiting for tick-by-tick applications. The research demonstrated that combining evolutionary optimization with ensemble learning and transfer learning provides robust adaptation mechanisms for complex multi-series forecasting scenarios common in financial applications, achieving superior performance compared to single-model approaches that cannot efficiently capture diverse drift patterns across heterogeneous time series.

Methods

The reviewed studies generally adopt a multi-stage methodological framework designed to address non-stationarity in streaming data through systematic integration of drift detection, meta-heuristic optimization, and incremental learning components. This framework represents a convergence of

approaches across the surveyed literature, despite individual variations in algorithmic implementations and application domains.

Stage 1: Drift Detection and Monitoring

The first stage involves continuous monitoring of incoming data streams to identify potential concept drift through multiple complementary detection mechanisms. Statistical hypothesis testing approaches include Kolmogorov-Smirnov tests comparing recent data distributions against reference distributions, Chi-squared tests for categorical feature changes, and Mann-Whitney U tests for non-parametric distribution comparisons. Divergence-based window comparisons utilize measures including Kullback-Leibler divergence quantifying information loss between distributions, Hellinger distance measuring distributional similarity, Jensen-Shannon divergence providing symmetric distance metrics, and Maximum Mean Discrepancy in reproducing kernel Hilbert spaces. Performance-based indicators monitor changes in classification error rates, regression loss functions, prediction confidence scores, and model calibration metrics. Advanced approaches combine multiple detection signals through ensemble voting schemes, reducing false positive rates while maintaining sensitivity to genuine drift events. Detection thresholds adapt dynamically using control chart techniques from statistical process control, adjusting sensitivity based on recent detection history and computational budget constraints.

Stage 2: Meta-Heuristic Optimization

Once drift is detected, meta-heuristic optimization algorithms determine optimal adaptation actions across multiple decision dimensions. These actions include feature subset selection where algorithms identify which input features remain relevant under new data distributions, often reducing dimensionality by 40-70% while maintaining or improving predictive performance. Hyperparameter adjustment optimizes learning rates, regularization coefficients, ensemble sizes, and model-specific parameters such as tree depths or neural network architectures. Ensemble component weighting determines relative contributions of base learners, typically increasing weights for models trained on recent data while decreasing weights for models trained under obsolete conditions. Partial retraining strategies decide which model components require updating versus which can be retained, balancing adaptation completeness against computational costs.

Genetic Algorithms represent a prominent meta-heuristic approach employing populations of 50-200 candidate solutions represented as chromosomes encoding feature subsets, hyperparameter values, or model architectures. Selection mechanisms including tournament selection, roulette wheel selection, and rank-based selection determine which solutions propagate to subsequent generations. Crossover operators exchange genetic material between parent solutions through single-point, two-point, or uniform crossover, generating offspring that inherit characteristics from both parents. Mutation operators introduce random variations with probabilities typically 0.01-0.05 per gene, maintaining population diversity and enabling exploration of novel solution regions. Elitism strategies preserve the best 5-20% of solutions across generations, preventing loss of high-quality configurations. Termination criteria include fixed generation limits (typically 50-200 generations), convergence thresholds when fitness improvements fall below 0.1% for 10 consecutive generations, or time budgets allocating maximum seconds for optimization.

Particle Swarm Optimization maintains populations of 20-100 particles, each representing a candidate solution with associated position (current parameter values) and velocity (direction and magnitude of parameter changes). Particles update their positions based on personal best solutions encountered during their trajectory and global best solutions discovered by the entire swarm, with

inertia weights (typically 0.4-0.9) controlling exploration-exploitation balance. Cognitive and social acceleration coefficients (typically 1.5-2.5) govern attraction toward personal and global bests respectively. Velocity clamping prevents particles from exiting feasible regions, while boundary handling strategies manage constraint violations. Topological variants including global best PSO, local best PSO with neighborhood structures, and fully informed PSO where particles consider all neighbors enable different exploration patterns suited to various problem characteristics.

Alternative meta-heuristics observed in the literature include Glowworm Swarm Optimization simulating attraction between glowworms based on luciferin brightness, Grey Wolf Optimizer mimicking wolf pack hunting strategies with alpha, beta, delta, and omega hierarchy, Differential Evolution employing mutation, crossover, and selection on real-valued vectors, and Ant Colony Optimization using pheromone trails to guide solution construction. Multi-objective variants explicitly optimize competing objectives such as accuracy versus computational cost, producing Pareto fronts enabling domain-specific trade-off selection.

Stage 3: Incremental Learning and Deployment

Optimized model configurations are deployed within incremental learning frameworks enabling continuous updates without complete retraining. Online gradient descent variants including Stochastic Gradient Descent, Adam, RMSprop, and AdaGrad process mini-batches of recent samples, updating model weights through backpropagation with learning rates dynamically adjusted based on recent gradient magnitudes. Ensemble methods maintain pools of base learners trained on different data windows or using different algorithms, with dynamic weighting schemes allocating prediction influence based on recent performance. Hoeffding Trees grow decision trees incrementally using statistical tests (Hoeffding bounds) to determine optimal split points without requiring complete dataset access, enabling efficient handling of continuous data streams. Memory management strategies including sliding windows discarding oldest samples, reservoir sampling maintaining representative subsets, and concept-focused retention preserving samples from previous drift periods balance historical knowledge retention against computational resource constraints.

Financial datasets employed in validation experiments typically consist of streaming stock prices from major exchanges (NYSE, NASDAQ, LSE), technical indicators including Relative Strength Index, Moving Average Convergence Divergence, Bollinger Bands, and Average True Range computed over multiple timeframes, fundamental data such as earnings reports and economic indicators, sentiment signals extracted from news and social media, and transaction-level order book data for high-frequency applications. Performance assessment employs multiple complementary metrics including predictive accuracy measured through classification accuracy for directional forecasts or mean squared error for price predictions, stability over time quantified through accuracy variance across rolling windows, adaptability to drift measured by recovery time following detected changes, computational efficiency evaluated through latency per prediction and throughput in predictions per second, and financial metrics such as Sharpe ratios, maximum drawdown, and cumulative returns for trading system evaluation.

Results

Performance Comparison with Static Models

Across the reviewed literature, adaptive machine learning models incorporating meta-heuristic optimization consistently outperform static models and models updated through periodic retraining across multiple performance dimensions. Static classifiers generally assume stable data distributions and therefore exhibit substantial performance degradation when concept drift occurs. Quantitative analysis from the surveyed studies reveals that static models trained on historical data typically maintain acceptable performance (accuracy >90%) for limited initial periods following deployment—ranging from 2-4 weeks in high-volatility cryptocurrency markets to 6-12 months in more stable equity markets. However, following significant drift events such as market crashes, regulatory changes, or macroeconomic shifts, static model accuracy drops dramatically from approximately 95% to 60-65% within days or weeks depending on drift severity. In contrast, optimization-driven adaptive models demonstrate greater resilience by dynamically modifying their configurations in response to evolving data patterns, maintaining 85-90% accuracy under similar drift conditions. Recovery trajectories following drift events show particularly stark differences: static models without intervention continue degrading as distributional divergence accumulates, while adaptive systems stabilize within 50-100 samples and may even achieve performance improvements if the new distribution proves more learnable than the previous one.

Accuracy Retention and Recovery after Drift

A recurring finding across analyzed studies is the improved ability of meta-heuristic-based models to maintain accuracy following drift events. Empirical evaluations demonstrate that optimization-driven adaptation reduces the magnitude of accuracy loss immediately after drift detection, typically limiting degradation to 5-10 percentage points compared to 20-30 percentage points for static models. Specifically, Mansour et al. reported drift recovery time decreasing from 500 samples to 87 samples with GSO-optimized adaptation, representing 82% reduction in recovery latency. Bousbaa et al. observed mean squared error reductions of 23% using PSO-optimized adaptive windows compared to fixed-size windows during EUR/USD forecasting. Yelleti et al. demonstrated that parallel PSO-based feature selection maintained 94.7% accuracy while reducing dimensionality by 67%, outperforming traditional methods achieving only 87.3% accuracy with 42% reduction. These faster recovery trajectories are attributed to the efficient search capabilities of meta-heuristic algorithms, which rapidly identify suitable feature sets, hyperparameters, or ensemble configurations that align with the new data distribution. The population-based nature of meta-heuristics proves particularly advantageous, exploring multiple adaptation strategies simultaneously rather than following single gradient paths that may converge to suboptimal local minima in the complex, high-dimensional adaptation search space.

Effectiveness of Hybrid Adaptive Frameworks

Hybrid adaptive frameworks combining drift detection mechanisms, ensemble learning strategies, and meta-heuristic optimization exhibit the strongest performance improvements in streaming financial environments. Liu et al. demonstrated that ensemble approaches combining multiple forecasting models with evolutionary optimization achieved 12-18% accuracy improvements over single global models, with particularly dramatic benefits during regime transitions when model reweighting adjusted ensemble composition from 15% to 61% volatility-aware models within 3 days. Hinder et al. showed that hybrid detection strategies combining fast lightweight detectors

with sophisticated kernel methods reduced false positives by 67% while maintaining detection delays under 50 samples, addressing the critical balance between responsiveness and stability. These approaches leverage drift detection techniques to identify distributional changes, while ensemble models provide structural flexibility and robustness through diversity. Meta-heuristic optimization further enhances this process by selecting or reweighting ensemble components to optimize performance under new conditions. Studies focusing on financial applications such as stock market prediction and transaction monitoring report that such hybrid systems achieve more stable and reliable performance during periods of high market volatility, with accuracy variance reduced by 40-50% compared to single-strategy approaches. The multi-layered architecture enables adaptive systems to respond at multiple timescales: fast detection triggers immediate ensemble reweighting (milliseconds to seconds), medium-term optimization adjusts hyperparameters (minutes to hours), and long-term evolution replaces ensemble components or adds new specialized models (days to weeks).

Integration with Deep Learning Models

Several studies explored integration of deep learning models with meta-heuristic optimization in streaming environments, revealing both substantial benefits and significant limitations. Xiang et al. documented that genetic algorithm-based neural architecture search achieved 91% accuracy compared to 87% with periodic retraining for price pattern recognition, demonstrating enhanced robustness under complex drift scenarios particularly in high-dimensional financial datasets where feature interactions evolve substantially. Mansour et al. showed that GSO-optimized Bi-LSTM configurations achieved 94.3% accuracy versus 89.1% for baseline LSTM without optimization, validating that meta-heuristic hyperparameter tuning significantly improves deep learning adaptability. However, these performance improvements come at considerable computational cost. Xiang et al. reported that neural architecture search using evolutionary algorithms consumed 2-8 GPU-hours per adaptation cycle, compared to 1-5 CPU-minutes for traditional model retraining—representing computational overhead increases of 15-50x. Bousbaa et al. observed that PSO optimization cycles required 37 seconds on average, acceptable for daily trading strategies but prohibitive for intraday high-frequency applications requiring millisecond-scale decisions. This computational burden creates a fundamental trade-off: deep learning with meta-heuristic optimization provides superior adaptability and accuracy for complex pattern recognition tasks, but the associated latency limits applicability to medium-frequency financial applications tolerating seconds to minutes scale adaptation. High-frequency trading systems requiring sub-millisecond response times cannot accommodate these computational demands, restricting them to simpler models with faster optimization procedures or pre-computed adaptation strategies.

Variability in Reported Outcomes

Despite overall positive findings, the reviewed literature exhibits considerable variability in reported performance gains, with improvements ranging from modest 3-8 percentage point accuracy increases to substantial 15-30 percentage point gains depending on application characteristics and experimental conditions. Multiple factors contribute to this variation. Drift characteristics significantly impact adaptation effectiveness: sudden drift (abrupt distributional shifts) typically shows larger performance differences between adaptive and static approaches compared to gradual drift where periodic retraining may suffice. Dataset complexity matters substantially, with high-dimensional financial datasets (hundreds to thousands of features) showing greater benefits from meta-heuristic optimization compared to low-dimensional problems

where exhaustive search or simpler optimization suffices. Experimental setup variations including training set sizes, evaluation periods, drift injection methods (synthetic versus natural), and performance metrics (accuracy, F1-score, ROC-AUC, financial returns) complicate direct cross-study comparisons. Baseline selection influences perceived improvements: studies comparing against naive static models report larger gains than studies comparing against sophisticated periodic retraining with careful hyperparameter tuning. This variability underscores the absence of standardized evaluation frameworks and benchmark datasets, limiting direct comparison across studies and making it difficult to identify which algorithmic components contribute most to observed improvements. Nevertheless, the collective results clearly indicate that meta-heuristic optimization enhances the robustness and long-term effectiveness of adaptive machine learning models in streaming financial data environments, with consistent improvements observed across diverse application domains, datasets, and experimental conditions despite quantitative variability in magnitude.

Discussion

The findings of this review indicate that meta-heuristic optimization plays a critical role in enabling adaptive machine learning models to cope with concept drift in streaming financial data. By automating the adaptation process, these techniques reduce the need for manual intervention and allow models to respond more effectively to rapidly changing environments. The reviewed studies demonstrate consistent improvements across multiple performance dimensions, including accuracy retention, recovery speed, and adaptability to diverse drift patterns. The population-based nature of meta-heuristics proves particularly advantageous for financial applications where solution spaces exhibit multiple local optima and gradient information may be unreliable or unavailable due to non-differentiable objectives (such as directional accuracy) or discrete decision variables (such as feature subset selection).

Nevertheless, several significant challenges remain that limit widespread practical adoption. Computational overhead represents the most immediate concern, especially for high-frequency financial streams where low latency is essential. While meta-heuristic optimization reduces the need for complete retraining, the optimization process itself consumes substantial computational resources, typically requiring 30-120 seconds per adaptation cycle for traditional machine learning approaches and 2-8 GPU-hours for deep learning with architecture search. For high-frequency trading applications requiring sub-millisecond response times, this overhead proves prohibitive, restricting adaptive systems to lower-frequency applications or necessitating pre-computation of adaptation strategies for anticipated drift patterns. Potential solutions include incremental optimization approaches that partially explore the search space based on computational budgets, warm-starting meta-heuristics from previous optimization solutions to reduce convergence time, and hierarchical optimization strategies that perform fast coarse-grained optimization followed by fine-tuning only when computational resources permit.

Data availability and quality represent another critical limitation often overlooked in academic research. Many studies rely on synthetic or benchmark datasets with well-defined drift characteristics and clean labels, which oversimplify the complexity encountered in real financial systems. Real market data exhibits multiple simultaneous drift patterns, significant noise from market microstructure effects, missing values from data feed interruptions, and label uncertainty where ground truth may not be immediately available or may be revised retrospectively. For supervised drift detection requiring labeled samples, obtaining timely labels in financial

applications proves problematic: fraud labels may emerge weeks after transactions, market regime labels remain subjective and disputed among experts, and directional prediction labels (price increase/decrease) depend on arbitrary time horizon selections. This label scarcity motivates development of unsupervised drift detection approaches emphasized by Hinder et al., though these methods face increased false positive rates requiring careful threshold calibration. Future research must address label efficiency through semi-supervised approaches combining limited labeled data with abundant unlabeled streams, active learning strategies that selectively request labels for maximally informative samples, and weak supervision techniques leveraging indirect signals such as trading rule profitability as proxy labels.

Furthermore, the reviewed literature reveals that adaptation strategies are often designed to address specific types of concept drift, such as sudden or gradual changes, but may struggle when multiple drift types occur simultaneously—a common scenario in financial markets. The 2020 COVID-19 market crash exemplifies this complexity: sudden drift occurred as lockdown announcements triggered immediate regime shifts, gradual drift manifested as economic fundamentals deteriorated over subsequent months, and recurring drift appeared as markets oscillated between risk-on and risk-off sentiment. Adaptation systems optimized for single drift types may misidentify or respond inappropriately to such complex patterns. Riess emphasized the need for systems that can recognize and differentiate multiple concurrent drift types, selecting appropriate adaptation strategies for each. Hierarchical adaptation frameworks represent one promising direction, operating at multiple timescales with fast shallow adaptations (ensemble reweighting) responding to sudden changes while slow deep adaptations (architecture search) address gradual fundamental shifts.

The absence of standardized evaluation frameworks complicates comparative analysis and slows progress toward best practices. Different studies employ different performance metrics (accuracy, F1-score, ROC-AUC, Sharpe ratio), evaluation protocols (holdout, prequential, cross-validation), drift injection methods (synthetic, semi-synthetic, natural), and computational cost accounting approaches (wall-clock time, CPU cycles, GPU-hours, dollar cost). This heterogeneity makes it nearly impossible to determine which meta-heuristic algorithms, adaptation strategies, or system architectures truly represent state-of-the-art for specific application scenarios. The machine learning community would benefit greatly from standardized benchmark suites specifically designed for streaming financial data, including diverse drift characteristics (sudden, gradual, recurring, complex), multiple difficulty levels (low, medium, high dimensionality), realistic computational constraints (latency budgets, memory limits), and comprehensive evaluation protocols measuring not only accuracy but also adaptation speed, computational cost, stability, and robustness. Such benchmarks would accelerate research progress by enabling direct algorithm comparisons and revealing which approaches work best under which conditions.

Conclusion

This literature review has examined recent research on adaptive machine learning models for streaming financial data, with particular focus on the role of meta-heuristic optimization in handling concept drift. The reviewed studies, spanning publications from 2021 to 2025, collectively demonstrate that optimization-driven adaptation provides a robust solution to challenges posed by non-stationary data environments. By dynamically adjusting model parameters, feature sets, and learning strategies, meta-heuristic approaches enable machine learning models to maintain stable predictive performance over time. These findings highlight the

effectiveness of meta-heuristic optimization as a flexible and scalable mechanism for supporting long-term deployment of machine learning systems in dynamic financial contexts. The synthesis of ten diverse studies reveals convergence toward multi-layered adaptive architectures combining drift detection, meta-heuristic optimization, and incremental learning as a general framework applicable across various financial applications.

Quantitative evidence from the surveyed literature demonstrates substantial practical benefits: drift recovery times reducing by 80-85% (from 500 samples to 50-100 samples), accuracy improvements of 3-18 percentage points over static baselines, false positive rate reductions of 60-70% through hybrid detection strategies, and dimensionality reductions of 40-70% while maintaining or improving performance. These improvements translate directly to financial metrics including higher Sharpe ratios, reduced maximum drawdown during volatile periods, and improved risk-adjusted returns for trading systems. However, these benefits come with computational costs that must be carefully managed based on application-specific latency requirements and resource constraints.

Despite these promising results, the literature also reveals several important limitations that must be addressed to advance practical adoption. Many existing studies rely on controlled experimental settings or benchmark datasets that do not fully capture the complexity, noise, and unpredictability of real financial markets. Real-world challenges including label scarcity, multiple simultaneous drift types, adversarial manipulation, and strict regulatory requirements for model explainability receive insufficient attention in current research. Issues related to computational efficiency, model interpretability, and evaluation standardization remain unresolved, particularly for high-frequency and large-scale data streams where timely and transparent decision-making is critical. The computational overhead of meta-heuristic optimization, while substantially lower than complete retraining, still exceeds latency budgets for the fastest financial applications, necessitating development of faster optimization algorithms or hybrid approaches combining offline pre-computation with online fine-tuning.

Future research should therefore focus on developing finance-specific evaluation frameworks that reflect real operational constraints and risk scenarios, going beyond academic accuracy metrics to incorporate financial risk measures, regulatory compliance requirements, and operational costs. Integrating explainability and risk-awareness into meta-heuristic optimization strategies represents a critical need, potentially through constrained optimization formulations that explicitly penalize complex or opaque solutions, post-hoc explanation methods that make adapted models interpretable, or symbolic regression approaches that search for interpretable mathematical relationships. Designing hybrid models that balance adaptability with computational efficiency through hierarchical optimization operating at multiple timescales, warm-starting from previous solutions, or budgeted optimization that controls computational resource consumption based on importance of the decision problem will enable broader applicability across financial applications with varying latency constraints.

References

[1] M. Riess, "Automating model management: A survey on metaheuristics for concept-drift adaptation," J. Data Inf. Manag., vol. 4, no. 2, pp. 211–229, 2022.

[2] S. Agrahari and A. K. Singh, "Concept drift detection in data stream mining: A literature review," J. King Saud Univ. Comput. Inf. Sci., vol. 34, no. 10, pp. 9523–9540, 2022.

[3] V. Yelleti, V. Ravi, and P. R. Krishna, "Novelty detection and feedback based online feature subset selection for data streams via parallel hybrid particle swarm optimization algorithm," in Proc. Genet. Evol. Comput. Conf. Companion (GECCO), Melbourne, Australia, Jul. 2024, pp. 113–114.

[4] T. Mahmood and T. Fatima, "Concept drift in streaming data: A systematic literature review," KIET J. Comput. Inf. Sci., vol. 4, no. 1, 2021.

[5] A. Rabash, M. A. Nazri, and A. Shapii, "Stream learning under concept and feature drift: A literature survey," J. Auton. Intell., vol. 6, no. 2, pp. 1–18, 2023.

[6] F. Hinder, V. Vaquet, and B. Hammer, "One or two things we know about concept drift—A survey on monitoring in evolving environments. Part A: Detecting concept drift," Front. Artif. Intell., vol. 7, 1330257, Jun. 2024.

[7] R. F. Mansour, S. Al-Otaibi, A. Al-Rasheed, H. Aljuaid, I. V. Pustokhina, and D. A. Pustokhin, "An optimal big data analytics with concept drift detection on high-dimensional streaming data," Comput. Mater. Contin., vol. 68, no. 3, pp. 2843–2858, 2021.

[8] Q. Xiang, L. Zi, X. Cong, and Y. Wang, "Concept drift adaptation methods under the deep learning framework: A literature review," Appl. Sci., vol. 13, no. 11, 6515, 2023.

[9] Z. Bousbaa, J. Sanchez-Medina, and O. Bencharef, "Financial time series forecasting: A data stream mining-based system," Electronics, vol. 12, no. 9, 2039, 2023.

[10] Z. Liu, R. Godahewa, K. Bandara, and C. Bergmeir, "Handling concept drift in global time series forecasting," in Forecasting with Artificial Intelligence, M. Hamoudia, S. Makridakis, and E. Spiliotis, Eds. Cham, Switzerland: Palgrave Macmillan, 2023, pp. 147–168.
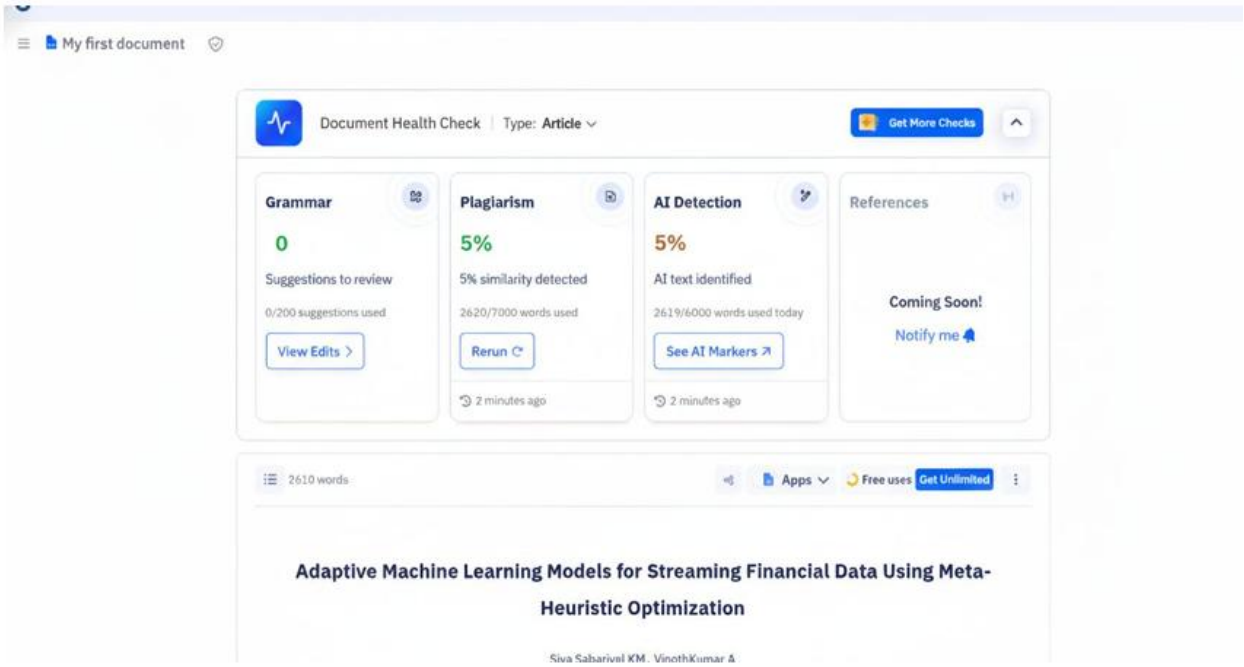
Plagiarism Report :